

XAI-Guided Interventional Retraining for Domain Generalization

2020-04-02

Suppose m is a model trained in one domain. Here is a simple method to obtain, from m , another model m' that hopefully generalizes better to other domains.

Let $\{h_u\}_{u=1}^F$ be all the feature units, and $\{z_c\}_{c=1}^C$ be the logit units. The weight between h_u and z_c in m is w_{uc} .

- For each training example (\mathbf{x}_i, y_i) , feed it to the model m and run an XAI method to obtain an heatmap e_i , which is what m considers as the core evidence in \mathbf{x}_i for the class y_i . Assume e_i is normalized to the interval $[0, 1]$.
- Create a “purified” input \mathbf{x}'_i by combining \mathbf{x}_i and e_i . The simplest way is to do pointwise multiplication. We can also convert e_i into a binary mask using a threshold before the pointwise multiplication.
- Feed \mathbf{x}'_i to m and compute the activation h'_{ui} of each feature unit h_u . Let $a_{ui} = h'_{ui}w_{uy_i}$. It is the contribution of feature h_u for classifying the purified example \mathbf{x}'_i into the class y_i .

Assume e_i correctly highlights the core evidence in \mathbf{x}_i for class y_i . Then the features h_u that do not focus on the core evidence would have low activation h'_{ui} on the “purified” example \mathbf{x}'_i , and hence a_{ui} would be low. Also, a_{ui} would be small when w_{uy_i} is small. As such, a_{ui} would be high only if: (1) the feature h_u focus on the core evidence for y_i and it is important for the class y_i according to the model m .

If a feature h_u is important for y_i according to m (i.e., with high w_{uy_i}), and it does not focus on the core evidence, it would lead to poor domain generalizability (DG). For better DG, we want to reduce the impact of such features. This can be achieved if we reduce the impact of all features with low a_{ui} . Note that by doing this we are also reducing the impact features with low w_{uy_i} , which should not bring about much impacts because the weights are low already.

- For each class c , define

$$A_{uc} = \frac{\sum_{i:y_i=c} a_{ui}}{\sum_{i:y_i=c} 1}$$

Note that $\sum_{i:y_i=c} 1$ is the number of examples in the class c in the training set. So, A_{uc} is the average of the a_{ui} ’s for the training examples in the class c .

Let say that the feature units h_u with $A_{uc} \geq \delta$ are *important* to the class c and the others are *unimportant* to c . Here, δ is a hyperparameter.

- To obtain m' ,
 - Set the weights for the features unimportant to the class c to 0, and keep them frozen during the following retraining process. The weights in the backbone are also frozen.
 - Re-initialize the weights for the important features, and retrain them.

This method feels more “heavy-handed” than previous methods. It is expected to lead to a decrease of the iid test error, and hopefully also lead to an increase of the ood test error.